


The background of the slide is a dense field of 3D-rendered numbers in various shades of blue. The numbers are of different sizes and are scattered across the frame, creating a sense of depth and complexity. Some numbers are in the foreground, while others are in the background, partially obscured.

# Spark Memory Management

TrendyTech

# Spark Memory Management

- ◆ Spark memory is divided into two parts
  - ◆ Heap memory : `spark.executor.memory` = 4GB
  - ◆ Off heap memory :  $\max(384, \text{spark.yarn.executor.memoryoverhead of } \text{spark.executor.memory})$   
=  $\max(384, 10\% \text{ of } 4\text{GB})$



# Spark Memory Management

- ◆ Heap memory : `spark.executor.memory` = 4GB
  - ◆ Reserved memory : 300 MB

|   |                        |  |
|---|------------------------|--|
| Off heap memory.<br>max(384MB, 10% of<br>4GB) | Java Heap memory - 4GB |  |
|   | 300MB<br>Reserved      |  |



# Spark Memory Management

- ◊ Heap memory : `spark.executor.memory` = 4GB
  - ◊ Reserved memory : 300 MB
  - ◊ Unified memory :  $(\text{Java Heap} - \text{Reserved Memory}) * \text{spark.memory.fraction}$ 
    - ◊ Unified Memory =  $(4\text{GB} - 300\text{MB}) * 0.6 = 3.70 \text{ GB} * 0.6 = 2.22 \text{ GB}$ .
  - ◊ User memory :  $(\text{Java Heap} - \text{Reserved Memory}) * (1 - \text{spark.memory.fraction})$ 
    - ◊ User memory =  $(4\text{GB} - 300\text{MB}) * (1 - 0.6) = 3.70 \text{ GB} * 0.4 = 1.48 \text{ GB}$ .

| Off heap memory.<br>max(384MB, 10% of<br>4GB) | Java Heap memory - 4GB |                           |  |
|---|------------------------|---------------------------|--|
|   | 300MB<br>Reserved      | User<br>Memory<br>1.48 GB | 2.22 GB =<br>Execution Memory + Storage Memory |

# Spark Memory Management

- ◆ Unified Memory is divided into
  - ◆ Execution memory :  $= (\text{Java Heap} - \text{Reserved Memory}) * \text{spark.memory.fraction} * (1.0 - \text{spark.memory.storageFraction})$ 
    - ◆ Execution memory =  $(4\text{GB} - 300\text{MB}) * 0.6 * (1-0.5) = 3.70 \text{ GB} * 0.6 * 0.5 = 1.11 \text{ GB}.$
  - ◆ Storage memory :  $= (\text{Java Heap} - \text{Reserved Memory}) * \text{spark.memory.fraction} * \text{spark.memory.storageFraction}$ 
    - ◆ Execution memory =  $(4\text{GB} - 300\text{MB}) * 0.6 * 0.5 = 3.70 \text{ GB} * 0.6 * 0.5 = 1.11 \text{ GB}.$

| Off heap memory.<br>max(384MB, 10% of<br>4GB) | Java Heap memory - 4GB |                           |                             |                           |
|---|------------------------|---------------------------|-----------------------------|---------------------------|
|   | 300MB<br>Reserved      | User<br>Memory<br>1.48 GB | Execution Memory<br>1.11 GB | Storage Memory<br>1.11 GB |

# Spark Memory Management

